

Suman Yadav

(318) 512-5501 | Monroe, LA | sumanydv615@gmail.com | linkedin.com/in/suman-ydv | github.com/Ydv-Suman

EDUCATION

University of Louisiana at Monroe

Bachelor of Science in Computer Science — GPA: 3.51/4.0

Monroe, LA

Expected May 2027

- **Relevant Coursework:** Artificial Intelligence, Data Structures & Algorithms, Object-Oriented Programming, Database Management Systems, Operating Systems, Advanced Discrete Structures, Calculus, Statistics & Probability

TECHNICAL SKILLS

- **Languages:** Python, Java, JavaScript, SQL, HTML5
- **Backend & Frameworks:** Spring Boot, Spring, Django, FastAPI, REST APIs, React, Tailwind CSS
- **Databases & Messaging:** MySQL, PostgreSQL, RabbitMQ, Kafka
- **Machine Learning & Data Science:** TensorFlow, scikit-learn, RAG, NumPy, Pandas, Matplotlib, Seaborn
- **DevOps & Tools:** AWS S3, Docker, Jupyter Notebook, Google Colab, Git, GitHub, Codex, Claude Code

EXPERIENCE

Startup Realm Technology

Web Developer Intern

Lalitpur, Nepal

May 2023 – Aug 2023

- Delivered 10+ responsive UI views using HTML, CSS, and JavaScript, contributing to a 10% increase in user engagement across multiple web surfaces.
- Engineered adaptive, cross-browser frontend components ensuring consistent UX across device and platform configurations, reducing reported rendering bugs during QA cycles.
- Collaborated in a 10-member Agile team maintaining code quality and version control via Git, enabling consistent sprint delivery through disciplined branching and peer code review.

HACKATHONS

Litmus | Python, FastAPI, React, PostgreSQL, AWS S3 | *GitHub* — Finalist

April 2026

- Engineered a candidate evaluation pipeline combining resume parsing, multi-source verification, and proctored video plus coding assessments; persisted candidate records and scoring results in PostgreSQL and stored resume and video artifacts in AWS S3.
- Implemented discrepancy detection between claimed experience and GitHub / LinkedIn activity to flag inflated profiles.
- Built a full-stack system with structured scoring and HR-facing insights; selected as a Hackathon Finalist.

PROJECTS

Atlas | Spring Boot, Java, PostgreSQL, Redis, Kafka, Docker | *GitHub*

June 2026 – Present

- Designing Atlas to connect to any backend repository (Spring Boot, FastAPI, Django, Express, NestJS), auto-detect its architecture and infrastructure dependencies, and expose an AI coding agent that proposes changes as reviewable diffs requiring static-analysis validation and human approval before merge.
- Completed auth-service end-to-end: OTP-based email verification backed by Redis, JWT issuance, and refresh token rotation, validated through full Postman test coverage before progressing to the next service.
- Architecting the platform as event-driven microservices (API Gateway, Auth, Repository, AI, Runtime, Deployment) communicating via Kafka, with planned multi-tenant security (JWT, RBAC, container isolation) and an observability stack.

Lernix | FastAPI, React, AWS S3, PostgreSQL, RAG, Docker | *GitHub*

Nov 2025 – Dec 2025

- Architected a cloud-native full-stack learning platform (FastAPI + React) on DigitalOcean; integrated AWS S3 pre-signed URL storage to offload all file I/O from the backend, enabling stateless scaling and reducing backend file-handling overhead by ~20%.
- Built an AI-powered document pipeline supporting chapter-level summarization, context-aware Q&A, and MCQ generation using chunking and embedding-based retrieval, structured as a pluggable RAG architecture.
- Implemented a learner analytics and recommendation layer analyzing content interaction patterns to deliver personalized study guidance, improving targeted content engagement by ~10–15%.

Trust AI | Python, FAISS, scikit-learn, RAG, FastAPI | *Live*

Dec 2025 – Jan 2026

- Built an LLM hallucination evaluation system using a 3-signal detection pipeline — retrieval overlap, semantic similarity scoring, and Natural Language Inference (NLI) contradiction analysis — grounded across 100+ document segments via FAISS vector search.
- Designed an interpretable trust scoring framework aggregating multi-signal outputs into a categorical risk score (Low / Medium / High) with human-readable explanations, enabling transparent and auditable AI assessments.
- Improved hallucination detection robustness across varied response lengths and claim densities by fusing independent signals into a composite reliability metric, reducing dependence on any single evaluation proxy.